Data Curation Techniques

Siddharth R

Syllabus

Unit 1: Introduction to Data Lifecycle
The data lifecycle (creation, storage, access, preservation) - Data objects and attribute types - Basic statistical descriptions of data - Measuring data similarity and dissimilarity - Database vs data warehouse - Data Quality: Why Preprocess the Data?

Unit 2: Data Preprocessing Techniques
Data cleaning workflow - Handling missing values, noisy data, outlier - Data integration : redundancy and correlation, duplication - Data transformation : normalization - data discretization using Binning, histogram, clusters - concept hierarchy generation for nominal data. Hands-on activity: Data preprocessing using Python libraries (Numpy, Pandas)

Unit 3: Data Reduction
Attribute subset selection method : forward selection, backward elimination, equal width and equal frequency histogram - sampling with and without replacement - data aggregation and summarization - overview of data cube - Basics of feature selection and feature extraction. Hands on: Data reduction using python (scikit-learn)

Unit 4: Data Management
ETL vs ELT - Data governance - data modeling and design - data integration and interoperability - Challenges of working with heterogeneous data sources - Common data formats (CSV, JSON, XML)- master data management - metadata - Discussion: Integration challenges in IoT data

Unit 5: Data Architecture
Principles of good data architecture - Architecture concepts : tight vs loose coupling - user access : single vs multi tenant - event driven architecture - data storage systems - storage abstraction - Hot, warm, and cold data - Discussion : Latest trends in storage using open source tools.

Text / Reference Books

1. *"Data Mining: Concepts and Techniques"* by Jiawei Han, Jian Pei, Hanghang Tong , Morgan Kaufmann, 2022, ISBN: 9780128117613 (For Unit 1 to 3)
2. "Fundamentals of Data Engineering" by Joe Reis and Matt Housley, O'Reilly Media, 2022, ISBN: ISBN: 9781098108304 (for Unit 4 and 5)
3. Latest related research articles from reputed journal/conferences

Assessment components

Quiz - 1 : 5%
Mid - Term : 25%
Quiz - 2 : 5 %
Lab Components : 25%
End Sem : 40%

Why Data Curation?

- The world is data rich but information poor.

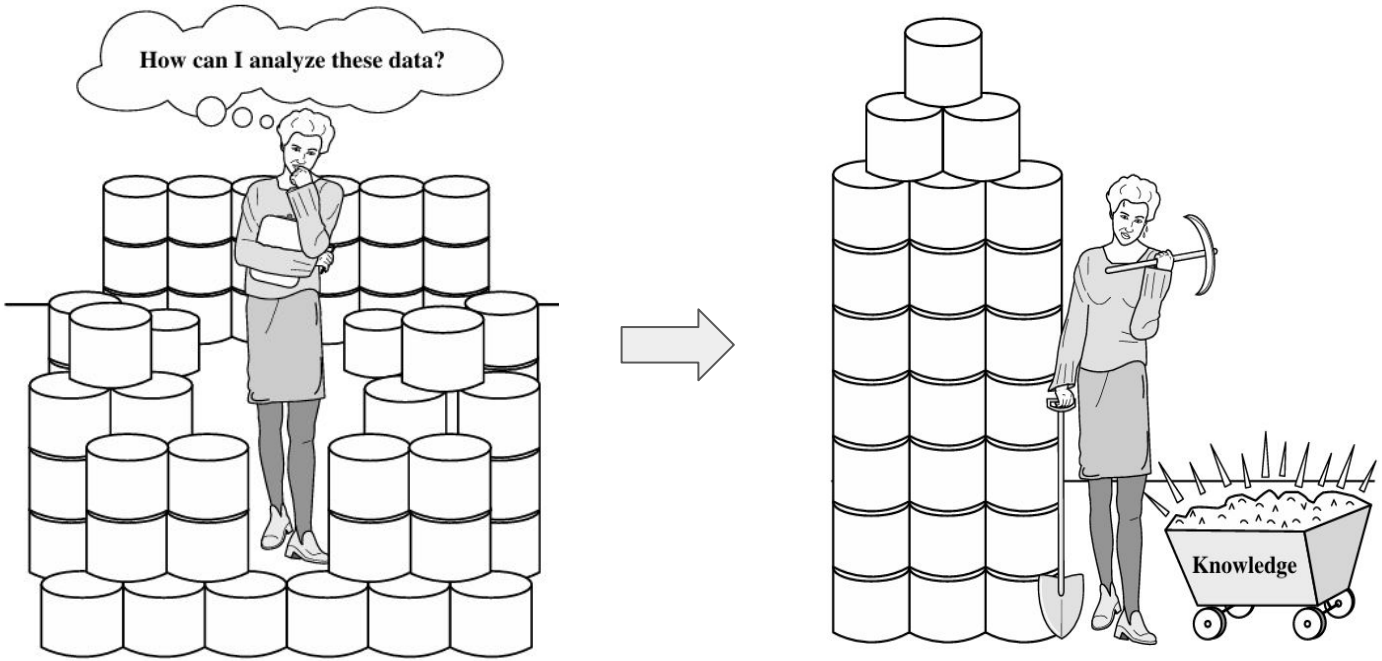

Why Data Curation?

Key steps:

- Data Cleaning
- Data Integration
- Data Selection
- Data Transformation

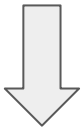

Data Analysis and Visualization

Data Life Cycle

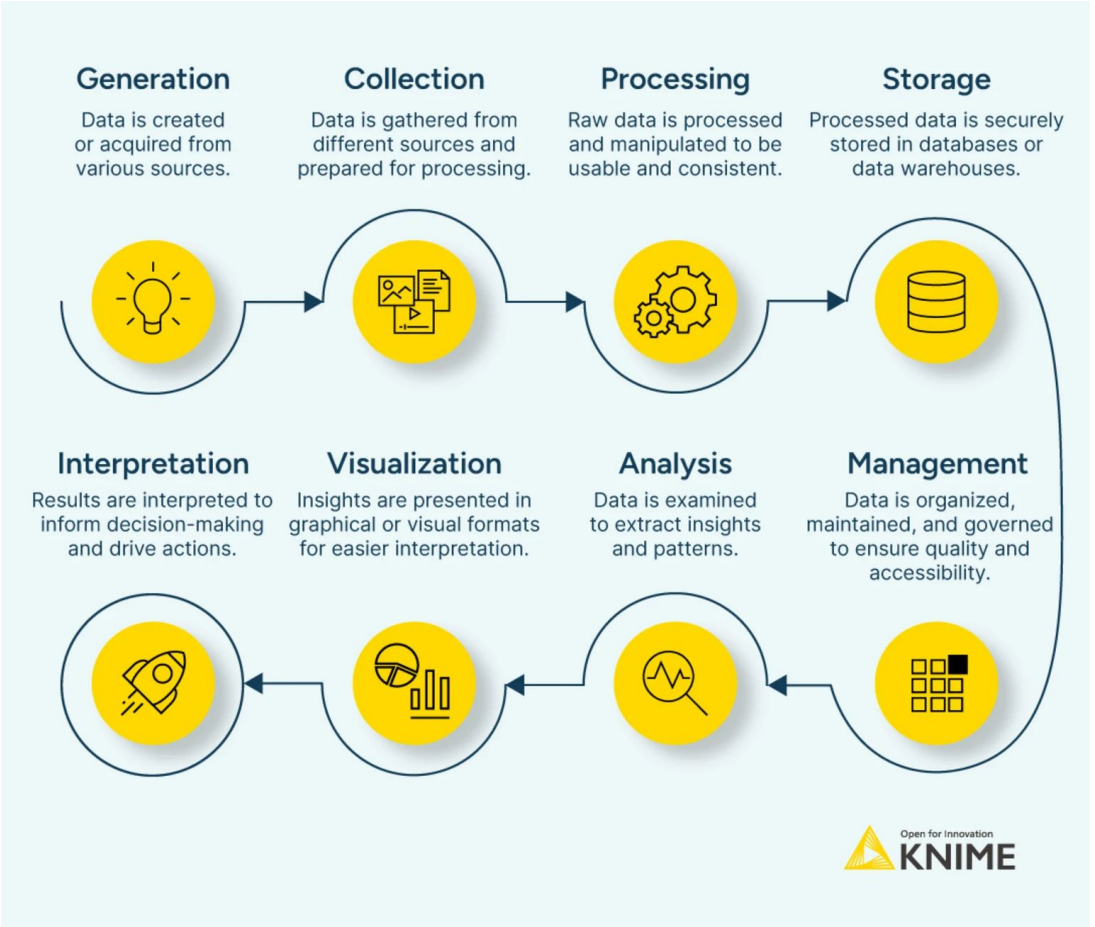

Source: https://www.knime.com/blog/the-data-lifecycle

Data Generation vs Data Collection

- Either active or passive
- Common sources:
  - Human-generated (e.g., surveys, forms)
  - Machine-generated (e.g., IoT sensors, logs)
  - Transactional systems (e.g., purchases, banking records)
  - Web scraping

What is Web Scraping (Common Data Collection)

- Extraction of data from website
- How the website content is presented?
- Is it structured ?
- How you are going to save the extracted content ?
- Can you name few common applications??
    - Sentiment Analysis?

Possible Ways to do Web Scraping

- Using Libraries in Programming Languages
- Browser Automation Tools
- APIs for Data Retrieval
- Headless Browsers
- No-Code or Low-Code Tools

Reading Assignment

Headless Browser - Puppeteer and Playwright

Key Applications :

1. E-commerce Price tracking
2. Social media monitoring

Common Data Sources

- Database data
  - A relational database is a collection of tables, each of which is assigned a unique name. Each table consists of a set of attributes (columns or fields) and usually stores a large set of tuples (records or rows).
  - Model : Entity - Relationship Model
- Data Warehouse
  - A data warehouse is a repository of information collected from multiple sources, stored under a unified schema, and usually residing at a single site.
  - Model : Data cube

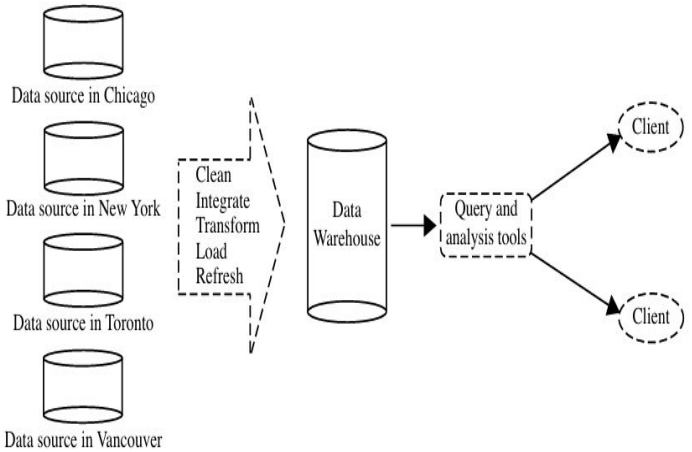

Data Objects

- Data sets are made up of data objects.
- A data object represents an entity
  - in a sales data, the objects may be customers, store items, and sales
  - In a university data ???
- Data objects can also be referred to as samples, examples, instances, data points, or objects.
- If the data objects are stored in a database, they are data tuples.
- Data objects are typically described by attributes.
- That is, the rows of a database correspond to the data objects, and the columns correspond to the attributes.
- The attribute, dimension, feature, and variable are often used interchangeably

Types of data attributes

- Nominal attribute:
  - Relating to name
  - Categorical
  - Data without any inherent order Examples: Gender, color, city
  - Is your aadhar number nominal?
  - How do you measure the central tendency ?

- Binary attribute:
  - Nominal attribute with only categories
  - Also called as Boolean
  - Give some examples ??
  - Symmetric vs Asymmetric binary attribute

Types of data attributes

- Ordinal attribute
  - Data with a specific order but without equal intervals between categories.
  - Examples: Education level, product rating (low, medium, high)
  - What is the preferred measure of central tendency?
- Whether the nominal, binary and ordinal data are quantitative or qualitative ?
- Numerical attribute:
  - Interval scaled attributes : No true zero point Example: Temperature in celsius / fahrenheit
  - Ratio scaled attributes : true zero point Example: Temperature in Kelvin
  - Continuous data: Data with infinite possible values within a given range. Examples: Height, weight, temperature
  - Discrete data: Data with a finite number of values. Examples: Number of children, number of products sold

Statistical Description of Data

- To have overall picture of your data
- Central tendency - measure the location of the middle or center of a data distribution.
- Dispersion of data - how are the data spread out
- Common approaches for Measuring the Central Tendency: Mean, Median, and Mode

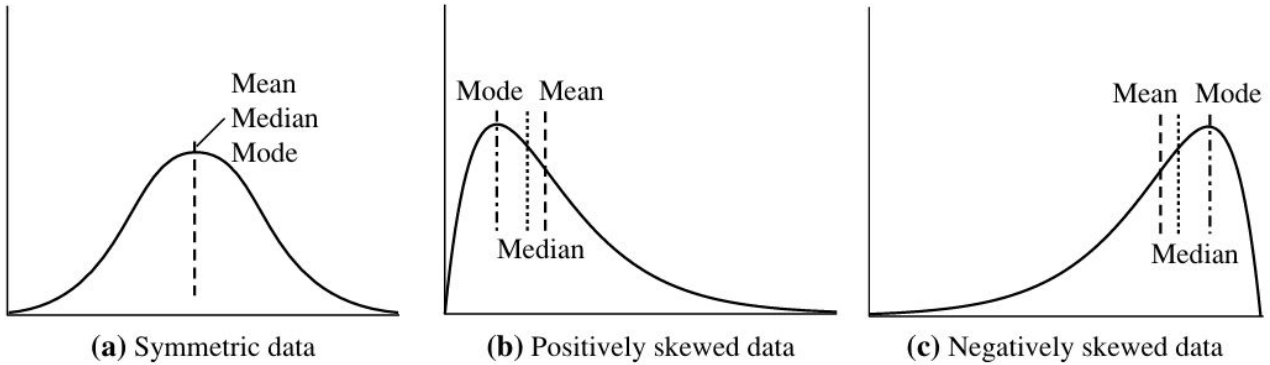

(a) Symmetric data **(b)** Positively skewed data **(c)** Negatively skewed data

Summary Statistics using Pandas [Refer Pandas code]

Key points to remember:

- Whereas size includes NaN values and count excludes the missing values.
- value_counts is a convenient shortcut to count the number of entries in each category of a variable.
- Split-Apply-Combine

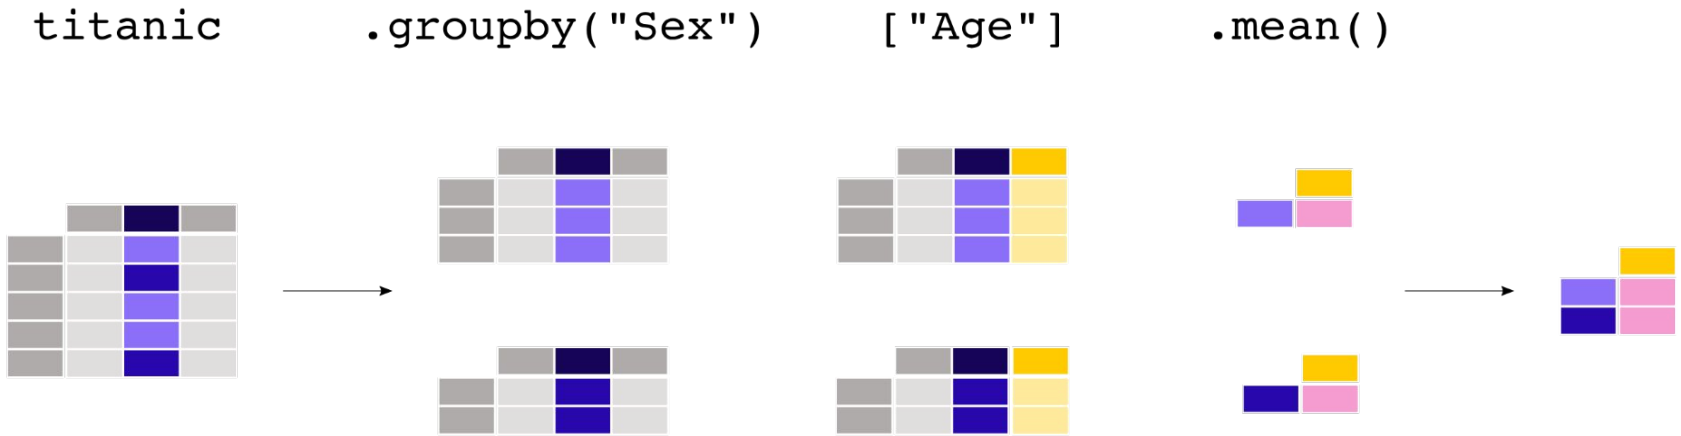

Measuring the Dispersion of Data

- Range : (Maximum - minimum)
- Standard Deviation : It provides a measure of the average distance between each data point and the mean
- Variance : The average of the squared differences between each data point and the mean.
- Quartiles: Quartiles divide the dataset into four equal parts when it is sorted.
    - Q1 (1st quartile): Median of the lower half (25th percentile).
    - Q2 (2nd quartile): Median of the dataset (50th percentile).
    - Q3 (3rd quartile): Median of the upper half (75th percentile).
- Interquartile Range (IQR) : The distance between the first and third quartiles , IQR = Q3 - Q1

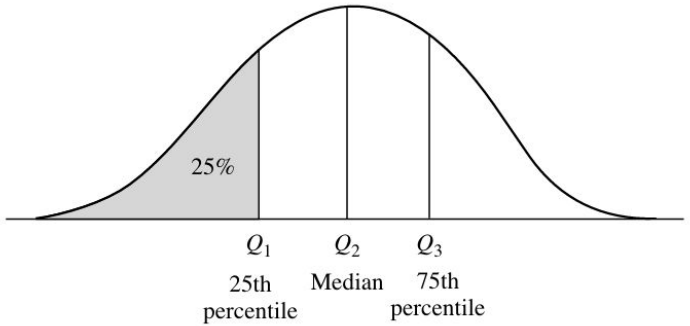

Measuring Data Similarity / Dissimilarity

- Why we need to measure the similarity / dissimilarity
- Also called as measure of proximity
- A similarity measure for two objects, i and j, will typically return the value 0 if the objects are unalike.
- The higher the similarity value, the greater the similarity between objects. (Typically,a value of 1 indicates complete similarity,that is,the objects are identical.)
- A Dissimilarity Measure works the opposite way.
- Data matrix vs Dissimilarity matrix

Data Matrix vs Dissimilarity Matrix

Data Matrix :This structure stores the n data objects in the form of a relational table, or n-by-p matrix (n objects X p attributes)

- Also called as “two-mode” matrix

$$\begin{bmatrix} x_{11} & \cdots & x_{1f} & \cdots & x_{1p} \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ x_{i1} & \cdots & x_{if} & \cdots & x_{ip} \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ x_{n1} & \cdots & x_{nf} & \cdots & x_{np} \end{bmatrix}.$$

Dissimilarity Matrix: This structure stores a collection of proximities that are available for all pairs of n objects.

where d(i, j) is the measured dissimilarity or difference between obj and j.

sim(i, j) = 1 - d(i, j), where *sim* is the similarity

- Also called as one-mode matrix

$$\begin{bmatrix} 0 & & & & \\ d(2,1) & 0 & & & \\ d(3,1) & d(3,2) & 0 & & \\ \vdots & \vdots & \vdots & & \\ d(n,1) & d(n,2) & \cdots & \cdots & 0 \end{bmatrix},$$

Measuring Dissimilarity - Nominal Attributes

Object Identifier	test-1 (nominal)
1	code A
2	code B
3	code C
4	code A

$$\begin{bmatrix} 0 & & & \\ d(2,1) & 0 & & \\ d(3,1) & d(3,2) & 0 & \\ d(4,1) & d(4,2) & d(4,3) & 0 \end{bmatrix}.$$

$$\begin{bmatrix} 0 & & & \\ 1 & 0 & & \\ 1 & 1 & 0 & \\ 0 & 1 & 1 & 0 \end{bmatrix}.$$

Euclidean Distance

- Measures the straight-line distance (as the crow flies).
- Diagonal movement allowed
- Euclidean distance is the shortest distance between any two points
- Mathematically, the Euclidean distance between the points x and y in two-dimensional plane is given by:

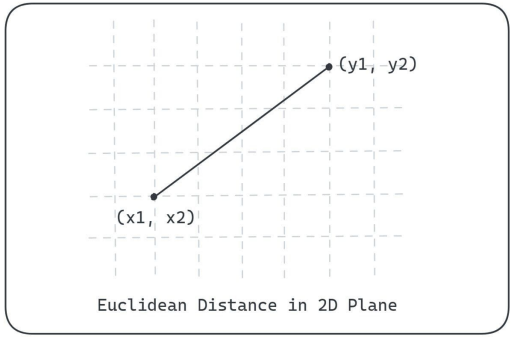

$$d(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2}$$

- Extending to n dimensions, the points x and y are of the form x = (x1, x2, …, xn) and y = (y1, y2, …, yn),

$$d(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{i=1}^{n} (x_i - y_i)^2}$$

Manhattan Distance (city-block)

- Measures the sum of the absolute differences of the coordinates.
- In a 2-D plane, the Manhattan distance between the points x and y is given by:

$$d(\mathbf{x}, \mathbf{y}) = |x_1 - y_1| + |x_2 - y_2|$$

- In n-dimensional space, where each point has n coordinates, the Manhattan distance is given by:

$$d(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^{n} |x_i - y_i|$$

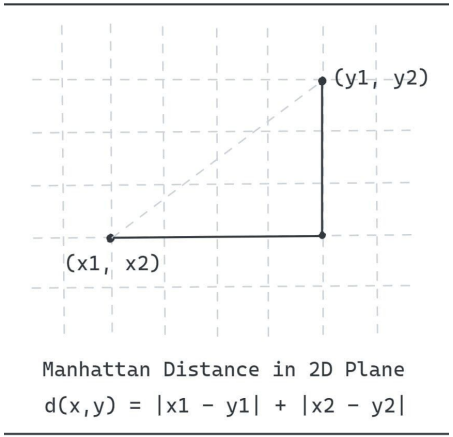

Minkowski Distance

- Minkowski distance is a generalization of the Euclidean and Manhattan distances. It is defined as

$$d(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^{n} |x_i - y_i|^p \right)^{1/p} \quad for\ p \geq 1$$

- If p=1, then Minkowski distance equation takes the same form as that of Manhattan distance (L1 - norm)
- Similarly, for p = 2, the Minkowski distance is equivalent to the Euclidean distance (L2-norm)

Chebyshev / Supremum distance

- It is a measure of distance that calculates the maximum difference along any coordinate dimension between two points in a multidimensional space.
- Also known as Lmax,L∞ norm or uniform norm

$$d_{\text{Chebyshev}}(P, Q) = \max_i(|p_i - q_i|)$$

- Let x_1 = (1, 2) and x_2 = (3, 5) , find Euclidean, Manhattan and Supremum

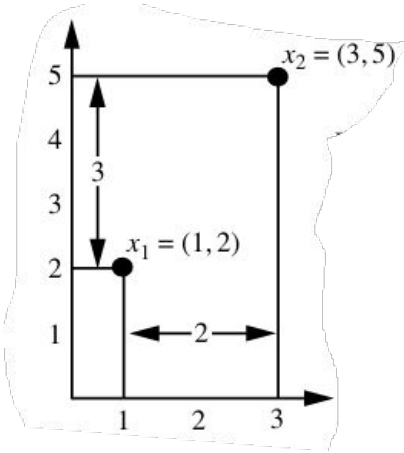

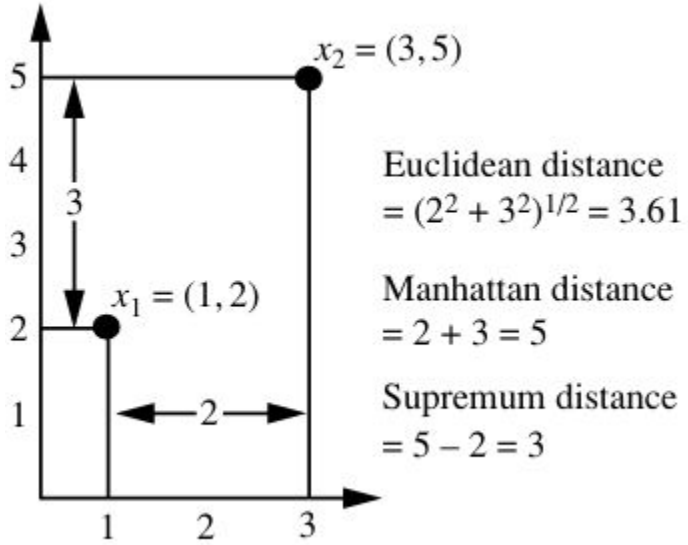
$x_2 = (3, 5)$
5
4
3
2
1
3
$x_1 = (1, 2)$
2
1
2
3
Euclidean distance
$= (2^2 + 3^2)^{1/2} = 3.61$
Manhattan distance
$= 2 + 3 = 5$
Supremum distance
$= 5 - 2 = 3$

Mathematical Properties

- Non-negativity : distance is always non-negative , $d(x, y) \geq 0$
- Identity of indiscernibles : distance between a point and itself is always 0. $d(x, y) = 0$ *if and only if* $x = y$
- Symmetry: The order of points doesn't matter in distance calculation. $d(x, y) = d(y, x)$
- Triangle inequality : The maximum difference between x and z along any dimension cannot be greater than the sum of the maximum differences from x to y and from y to z. $d(x, z) \leq d(x, y) + d(y, z)$

Jaccard Distance

- Quantify the dissimilarity between two sets of data.
- Derived from the Jaccard Index (or Similarity Coefficient)
- Jaccard Distance=1−Jaccard Index, where Jaccard Index=$|A \cap B|$ / $|A \cup B|$
- $A \cap B$: The number of elements common to both sets A and B (intersection).
- $A \cup B$: The number of unique elements in either set A or B (union).
- The Jaccard Index ranges from 0 to 1, where 1 indicates identical sets and 0 indicated disjoint sets
- Set A={1,2,3,4} and Set B={3,4,5,6} , $|A \cap B|$ =? , $|A \cup B|$ =? , Jaccard Index =?, Jaccard Distance =?
- Use case: Compare documents or text similarity, sets of pixels in images for similarity

Overview

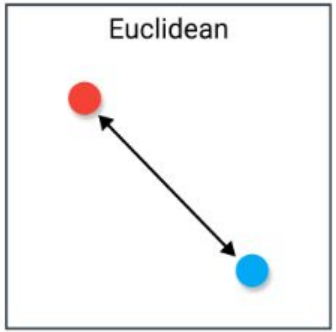
Euclidean

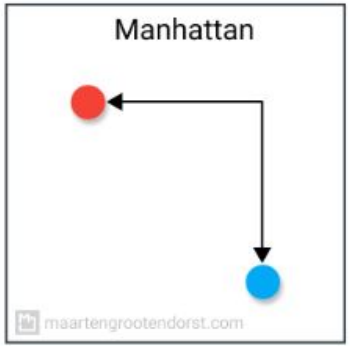
Manhattan
maartengrootendorst.com

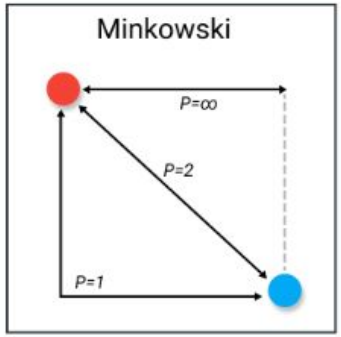
Minkowski
P=∞
P=2
P=1

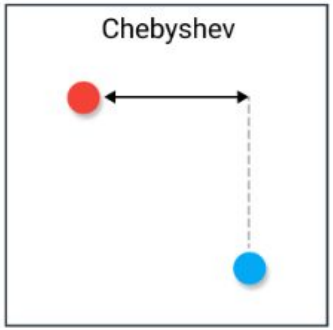
Chebyshev

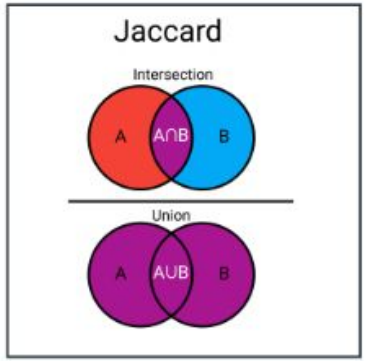
Jaccard
Intersection
A
A∩B
B
Union
A
A∪B
B

Thank You !!!